

# Getting started

## Description

A project implements the protocol "Read one copy, write all available copies" based on the ideas from "Site Recovery in Replicated Distributed Database Systems", 6th IEEE Distributed Computing Systems Conference, May 1986. User transactions can read and write logical objects at any available site, changed values replicate among available sites and they are marked in fail-lock tables for further recovery of unavailable sites. If at least one site is always up, it is possible to perform operations on available copies and failed sites can recover successfully after crashing.

The project consists from 2 java modules. The first one simulates 3 sites which communicate with each other via java RMI. The second one is used as a simulator for demonstration, it forces sites to fail or wake up and simulates user queries to the sites.

## Requirements

- Java version 1.8 or higher.

## Compile

### Windows and Linux

1. Compile from the Site project directory (Site/):

```
javac -d bin src/*.java
```

2. Compile from the SiteController project directory (SiteController/):

```
javac -d bin src/*.java
```

## Start

### Windows

1. Set Site/bin/ as a working directory;

2. Start RMI:

```
rmiregistry
```

3. Starting the application it is necessary to set option "-Djava.rmi.server.codebase" to the full path of the directory with binary files of the project. For example, "D:\Eclipse\workspace\Site\bin\".

To start single container for 3 sites of Site project:

```
java -Djava.rmi.server.codebase=file:D:\Eclipse\workspace\Site\bin\bin\ SiteImpl
```

OR start 3 sites separately:

```
java -Djava.rmi.server.codebase=file:D:\Eclipse\workspace\Site\bin\bin\ Site0
```

```
java -Djava.rmi.server.codebase=file:D:\Eclipse\workspace\Site\bin\bin\ Site1
```

```
java -Djava.rmi.server.codebase=file:D:\Eclipse\workspace\Site\bin\bin\ Site2
```

4. Set SiteController/bin/ as a working directory;

5. Start SiteController:

```
java SiteController
```

### Linux

From the command line:

1. Set Site/bin/ as a working directory;

2. Start RMI:

## rmiregistry

- Starting the application it is necessary to set option “-Djava.rmi.server.codebase” to the full path of the directory with binary files of the project. For example, “/tkuznets/workspace/Site/bin/”.

To start single container for 3 sites of Site project:

```
java -Djava.rmi.server.codebase=file:/tkuznets/workspace/Site/bin/ SiteImpl
```

OR start 3 sites separately:

```
java -Djava.rmi.server.codebase=file:/tkuznets/workspace/Site/bin/ Site0
```

```
java -Djava.rmi.server.codebase=file:/tkuznets/workspace/Site/bin/ Site1
```

```
java -Djava.rmi.server.codebase=file:/tkuznets/workspace/Site/bin/ Site2
```

- Set SiteController/bin/ as a working directory;

- Start SiteController:

```
java SiteController
```

## Commands

Commands for site controller:

| Command       | Format                               | Description                                |
|---------------|--------------------------------------|--------------------------------------------|
| HELP          | HELP                                 | Print help                                 |
| READ          | <site_num> READ <item>*              | Read data item at the site                 |
| WRITE         | <site_num> WRITE <item>*<br/><value> | Write data item at the site                |
| FAIL          | <site_num> FAIL                      | Simulate site failure                      |
| WAKE_UP       | <site_num> WAKE_UP                   | Simulate site recovery                     |
| PRINT_ITEMS   | <site_num> PRINT_ITEMS               | Print all values of data items of the site |
| PRINT_VECTOR  | <site_num> PRINT_VECTOR              | Print session vector of the site           |
| PRINT_FLOCKS  | <site_num> PRINT_FLOCKS              | Print all fail-locks stored at the site    |
| PRINT_UNREAD  | <site_num> PRINT_UNREAD              | Print unreadable data items                |
| PRINT_SESSION | <site_num> PRINT_SESSION             | Print scutal sesision number               |
| TIMER         | <site_num> TIMER true false          | Turn timer on/off                          |
| LOG           | <site_num> LOG true false            | Turn logging on/off                        |
| SLEEP         | <site_num> SLEEP true false          | Turn sleeping mode on/off                  |
| HELP          | HELP                                 | Print help message                         |
| Q             | Q                                    | Exit site controller                       |

\* item in range {X, Y, Z, L, M, N, A, B, C, D, E, F}

## Test scenarios

### Test 1

Simple test which shows read and write operations and replication of new values.

Input in the SiteController command line:

- >>> 0 PRINT\_ITEMS
- >>> 1 PRINT\_ITEMS
- >>> 2 PRINT\_ITEMS
- >>> 0 READ X
- >>> 0 WRITE X 0

6. >>> 0 READ X
7. >>> 1 READ X
8. >>> 2 READ X
9. >>> 0 PRINT\_ITEMS
10. >>> 1 PRINT\_ITEMS
11. >>> 2 WRITE Y 0
12. >>> 2 PRINT\_ITEMS
13. >>> 0 PRINT\_ITEMS

### Test 2

This test shows failure and recover of 2 sites and writing values while 1 or 2 sites are down.

Input in the SiteController command line:

1. >>> 0 FAIL
2. >>> 1 PRINT\_VECTOR
3. >>> 2 PRINT\_VECTOR
4. >>> 1 WRITE X 2
5. >>> 1 PRINT\_VECTOR
6. >>> 2 PRINT\_VECTOR
7. >>> 0 WAKE\_UP
8. >>> 0 PRINT\_ITEMS
9. >>> 1 PRINT\_VECTOR

### Test 3

This test shows the recovery process when one site fails after another and they recover in the same order.

Input in the SiteController command line:

1. >>> 0 FAIL
2. >>> 1 WRITE Z 0
3. >>> 1 FAIL
4. >>> 2 PRINT\_VECTOR
5. >>> 0 WAKE\_UP
6. >>> 0 PRINT\_ITEMS
7. >>> 0 PRINT\_VECTOR
8. >>> 2 PRINT\_VECTOR
9. >>> 2 WRITE Y 0
10. >>> 1 WAKE\_UP
11. >>> 1 PRINT\_ITEMS

## Test 4

This test shows the maintenance and collection of fail-locks.

Input in the SiteController command line:

1. >>> 0 FAIL
2. >>> 1 PRINT\_FLOCKS
3. >>> 2 PRINT\_FLOCKS
4. >>> 1 WRITE Z 0
5. >>> 1 PRINT\_FLOCKS
6. >>> 2 PRINT\_FLOCKS
7. >>> 2 WRITE A 0
8. >>> 1 PRINT\_FLOCKS
9. >>> 1 PRINT\_ITEMS
10. >>> 0 WAKE\_UP
11. >>> 0 PRINT\_ITEMS

## Test 5

This test shows the performance of the recovery if during failure none or all logical objects were updated.

Input in the SiteController command line:

1. >>> 0 TIMER true
2. >>> 0 LOG false
3. >>> 0 FAIL
4. >>> 0 WAKE\_UP
5. >>> 0 FAIL
6. >>> 1 WRITE X 0
7. >>> 1 WRITE Y 0
8. >>> 1 WRITE Z 0
9. >>> 1 WRITE L 0
10. >>> 1 WRITE M 0
11. >>> 1 WRITE K 0
12. >>> 1 WRITE A 0
13. >>> 1 WRITE B 0
14. >>> 1 WRITE C 0
15. >>> 1 WRITE D 0
16. >>> 1 WRITE E 0
17. >>> 1 WRITE F 0

18. >>> 1 PRINT\_FLOCKS

19. >>> 0 WAKE\_UP

### Test 6

Shows the ability to receive write and replicate commands during recovery and perform them after it. To send those commands it is necessary to change the mode of the site that will recover and to start second SiteController.

Input in the first SiteController command line:

1. >>> 0 SLEEP true
2. >>> 0 FAIL
3. >>> 0 WAKE\_UP

Input in the second SiteController command line:

1. >>> 0 WRITE X 5
2. >>> 1 WRITE Y 6